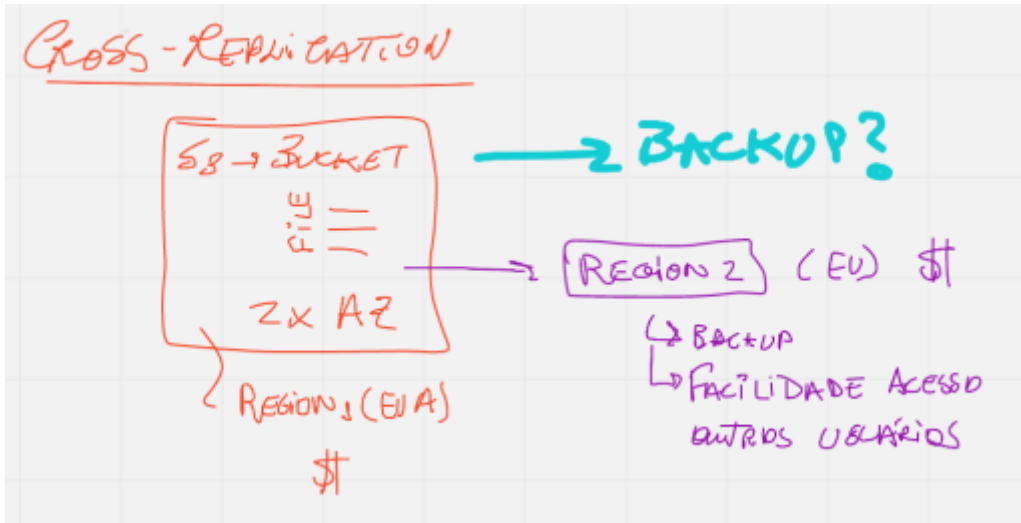
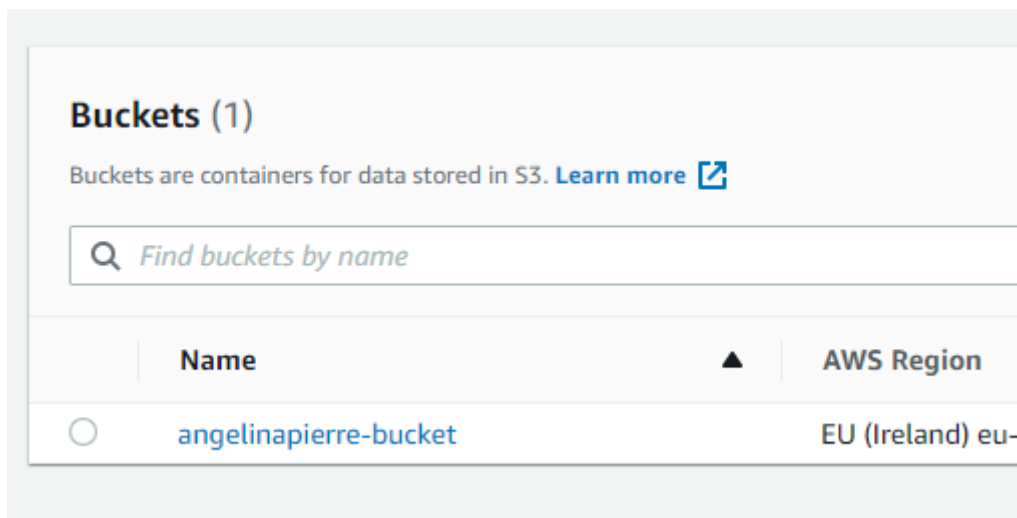


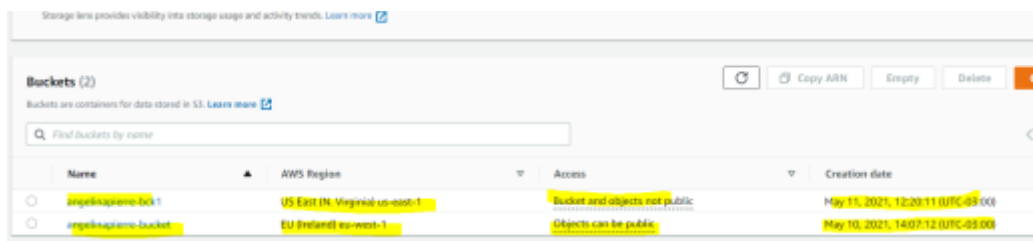
15 - HANDS ON : CONFIGURANDO CROSS REGION REPLICATION



- Nos temos uma bucket s3 onde será armazenado nossos arquivos.
- Aonde esta o backup desta bucket?
- Se utilizarmos a versão standard de armazenamento, de cara estaremos armazenando ele por default em duas availabilitys zones
- Queremos fazer algo mais, vamos armazenar em duas regioes diferentes.
 - Questoes de backup e facilidade de acesso aos usuarios.
- Podemos fazer com que todo o sistema seja replicado ou backup de dados para outra regioao.
- Temos que levar em consideração que tudo o que pagamos numa bucket iremos pagar na outra com a taxas da regioao.

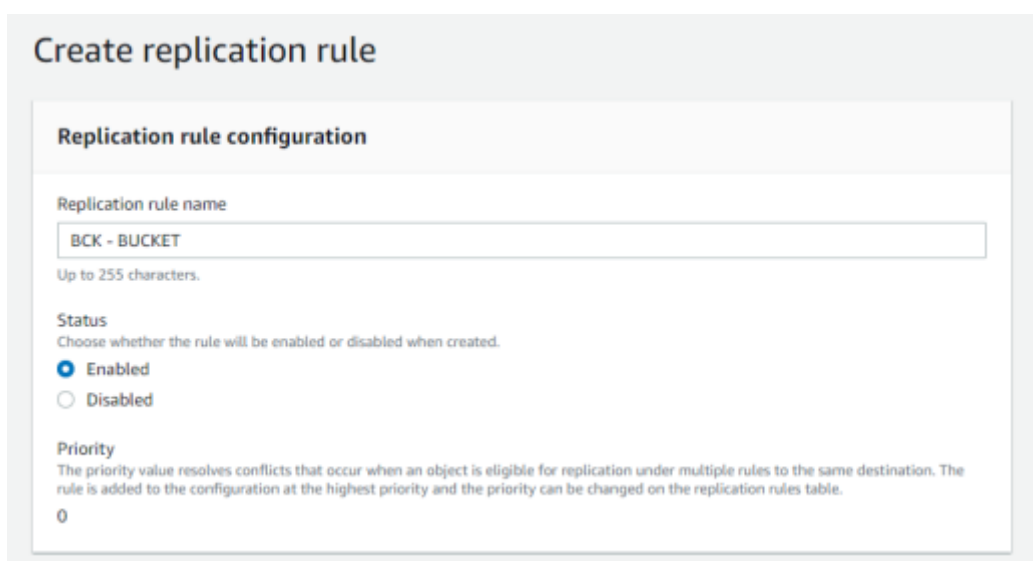


-Vamos criar outra bucket e deixa-la em outra regioao.



- Temos que apontar a bucket como sendo a replicaçãod e uma outra.

- Logo vamos no bucket que queremos fazer o backup e definimos as regras de replicação.



Source bucket

Source bucket name

angelinapierre-bucket

Source Region

EU (Ireland) eu-west-1

Choose a rule scope

☒ Limit the scope of this rule using one or more filters

☐ This rule applies to all objects in the bucket

Filter type

You can filter objects by prefix, object tags, or a combination of both.

Prefix

Add a filter to limit the scope of this rule to a single prefix.

Enter prefix

Don't include the bucket name in the prefix. Using certain characters in key names can cause problems with some applications and protocols.

Tags

You can limit the scope of this rule to the key value pairs added below.

Add tag

OU

Source bucket

Source bucket name

angelinapierre-bucket

Source Region

EU (Ireland) eu-west-1

Choose a rule scope

☐ Limit the scope of this rule using one or more filters

☒ This rule applies to all objects in the bucket

DEDSTINO

Destination

Destination

You can replicate objects across buckets in different AWS Regions (Cross-Region Replication) or you can replicate objects across buckets in the same AWS Region (Same-Region Replication). You can also specify a different bucket for each rule in the configuration. [Learn more](#) or see [Amazon S3 pricing](#)

☒ Choose a bucket in this account

☐ Specify a bucket in another account

Bucket name

Choose the bucket that will receive replicated objects.

angelinapierre-bkp

Browse S3

Destination Region

US East (N. Virginia) us-east-1

STORAGE CLASS RULES

Destination storage class

Amazon S3 offers a range of storage classes designed for different use cases. [Learn more](#) or see [Amazon S3 pricing](#)

☒ Change the storage class for the replicated objects

Storage class

	Storage class	Designed for	Availability Zones	Min storage duration	Price class
<input checked="" type="radio"/>	Standard	Frequently accessed data	≥ 3	-	-
<input type="radio"/>	Intelligent-Tiering	Long-lived data with changing or unknown access patterns	≥ 3	30 days	-
<input type="radio"/>	Standard-IA	Long-lived, infrequently accessed data	≥ 3	30 days	1
<input type="radio"/>	One Zone-IA	Long-lived, infrequently accessed, non-critical data	1	30 days	1
<input type="radio"/>	Glacier	Long-term data archiving with retrieval times ranging from minutes to hours	≥ 3	90 days	-
<input type="radio"/>	Glacier Deep Archive	Long-term data archiving with retrieval times within 12 hours	≥ 3	180 days	-
<input type="radio"/>	Reduced redundancy	Frequently accessed, non-critical data	≥ 3	-	-

resumo de acesso
BUCKET1 → S3-STANDARD(EUA)
BUCKET2 → S3-STANDARD(M.VIRGINIA)
(BACK) → USAR OUTRO
TIPO DE TIER

IAM role

IAM role

Create new role

View

Replication rules

Replication enables automatic and asynchronous copying of objects across buckets in the same or different AWS Regions. A replication configuration is a set of rules that define what actions should be applied to a group of objects during replication.

Replication configuration settings

Configuration settings affect all replication rules in the bucket

Source bucket

my-bucket-1

Source Region

us-east-1

IAM role

arn:aws:iam::123456789012:role/AmazonS3OutpostsReplicationRole

Replication rules (1)

Use replication rules to define options you want Amazon S3 to apply during replication, such as server-side encryption, replica ownership, transiting replicas to another storage class, and more. [Learn more](#)

View details

Edit rule

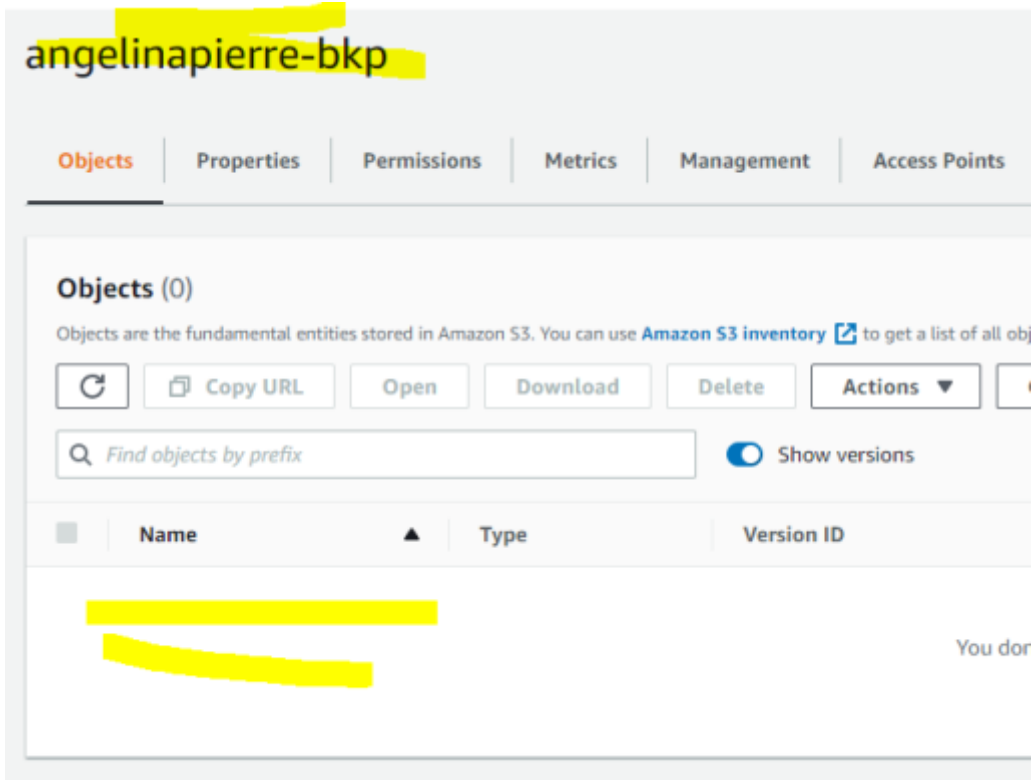
Delete

Actions

Create replication rule

Replication rule name	Status	Destination bucket	Destination Region	Priority	Scope	Storage class	Replica owner	Replication Time Control	KMS-encrypted objects	Replica modification ops
BUCKET1-TO-BUCKET2	Enabled	s3://my-bucket-2	us-east-1	1	Entire bucket	Standard	Same as source	Disabled	Do not replicate	Disabled

AGORA TEMOS QUE VERIFICAR SE OS ARQUIVOS FORAM REALMENTE REPLICADOS.



??????

- Quando voce realiza a sincronização da replicação, ele não copia os arquivos ja existentes , somente os futuros, teremos que forçar via linha de comando que ele copie os arquivos ja existentes de uma bucket para outra.

```
C:\Users\GABI>aws s3 ls
2021-05-11 12:24:37 angelinapierre-bkp
2021-05-11 12:32:28 angelinapierre-bucket
C:\Users\GABI>
```

comandod e copia

```
aws s3 cp -- recursive bucket_origem bucket_destino
```

```
C:\Users\GABI>aws s3 cp --recursive s3://angelinapierre-bucket s3://angelinapierre-bkp
copy: s3://angelinapierre-bucket/texto-1.txt to s3://angelinapierre-bkp/texto-1.txt
C:\Users\GABI>
```

angelinapierre-bkp

Objects

Properties


Permissions

Metrics

Ma

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon](#)



 Copy URL

Open

Download



 Find objects by prefix

<input type="checkbox"/>	Name	▲	Type
<input type="checkbox"/>	 texto-1.txt		txt